

# Time Series Analysis & Forecasting

## Class 9

Arnab Bose, Ph.D.

MSc Analytics

University of Chicago

# Cointegration

---

Regression of a non-stationary TS on another non-stationary TS

$$Y_t = \beta_0 + \beta_1 X_t + u_t$$

i.e.  $u_t = Y_t - \beta_0 - \beta_1 X_t$

where  $Y_t, X_t \sim I(d)$

$$u_t \sim I(0)$$

The linear combination cancels out the stochastic trends in  $X_t$  and  $Y_t$ .

$X_t$  and  $Y_t$  are said to be cointegrated

To check for cointegration, verify that the residuals  $u_t$  are  $I(0)$  or stationary

# Error Correction Model

---

From cointegration,  $X_t$  and  $Y_t$  have a long term equilibrium relationship.

But that may get off-balance in the short term

$$\Delta Y_t = \alpha_0 + \alpha_1 \Delta X_t + \alpha_2 u_{t-1} + \varepsilon_t$$

where  $\Delta$  is the first difference operator

$\varepsilon_t$  is the random error

$u_{t-1}$  one-period lagged error from cointegration regression

Note that  $\alpha_2 < 0$  to drive short term disequilibrium to equilibrium

# R code - Cointegration

---

```
library("quantmod", lib.loc=~R/win-library/3.3")
getSymbols("SPY", from="2013-01-01", to="2013-12-31")
getSymbols("IVV", from="2013-01-01", to="2013-12-31")
spyAdj <- unclass(SPY$SPY.Adjusted)
ivvAdj <- unclass(IVV$IVV.Adjusted)
adf.test(spyAdj)
adf.test(ivvAdj)
kpss.test(spyAdj, null="T")
kpss.test(ivvAdj, null="T")
ivv_spy <- lm(ivvAdj ~ spyAdj + 0)
regB <- lm(ewcAdj ~ ewaAdj + 0)
adf.test(ivv_spy$residuals)
adf.test(spy_ivv$residuals)
kpss.test(ivv_spy$residuals, null="T")
kpss.test(spy_ivv$residuals, null="T")
coef(ivv_spy)[1] * coef(spy_ivv)[1]#should be close to 1
```

# Granger Causality Test

---

Two TS  $X_t$  and  $Y_t$  - determine which one is useful in forecasting the other.

Use this test to determine G-causality

$H_0$ :  $X_t$  does not Granger-cause  $Y_t$

2 regressions

$$Y_t = a_0 + a_1 Y_{t-1} + \dots + a_m Y_{t-m} + \text{residual}$$

$$Y_t = a_0 + a_1 Y_{t-1} + \dots + a_m Y_{t-m} + b_p X_{t-p} + \dots + b_q X_{t-q}$$

Use F-test whose null hypotheses is no explanatory power is added by  $X_t$

Note if TS are cointegrated, then there must be a G-causality between them – either one-way or in both directions

# R code – Granger Causality Test

---

```
library("lmtest", lib.loc="~/R/win-library/3.3")  
grangertest(ChickEgg[, 1], ChickEgg[, 2], order = 3)
```

# Frequency Domain Representation

---

Frequency domain of a stationary TS representation

$$Y_t = \sum_{k=1}^T [a_k \sin(2\pi f_k t) + b_k \cos(2\pi f_k t)]$$

where

$$f_k = \frac{k}{T} \text{ } k \text{ is the \# of harmonics (Fourier frequencies)}$$

$$a_k = \frac{2}{T} \sum_{t=1}^T [\cos(2\pi f_k t)]$$

$$b_k = \frac{2}{T} \sum_{t=1}^T [\sin(2\pi f_k t)]$$

The auto-covariance is given by

$$\gamma_k = \sum_{j=1}^T \sigma_j^2 \cos(2\pi f_j t)$$

And the periodogram is given by

$$I(f_k) = \frac{T}{2} (a_k^2 + b_k^2)$$

# Frequency Domain Representation

---

The periodogram is given by

$$I(f_k) = \frac{T}{2} (a_k^2 + b_k^2)$$

- The periodogram is quickly computed using Fourier transform and is a “rough” estimate of the spectral density. Conversely, the periodogram is smoothed and scaled to produce the spectrum of the spectral density function.
- Generally if the frequency  $f_k = \frac{k}{T}$  is important (not), then  $I(f_k)$  will be large (small). The height of the periodogram shows the relative strength of sine-cosine pairs at various frequencies in the overall behavior of the TS.
- It can be shown that

$$\sum_{k=1}^T [I(f_k)] = \sigma^2$$



# R code – Frequency Domain Representation

---

```
library("forecast", lib.loc="~/R/win-library/3.3")
library("xts", lib.loc="~/R/win-library/3.3")
library("TSA", lib.loc="~/R/win-library/3.3")
data("USAccDeaths")
plot(as.xts(USAccDeaths), major.format = "%y-%m")
p <- periodogram(USAccDeaths)
p
max_freq <- p$freq[which.max(p$spec)]
seasonality <- 1/max_freq
seasonality

# white noise
periodogram(rnorm(1000))
```

# R code – Frequency Domain Representation

---

*# AR & MA models with positive and negative coefficients*

*par(mfrow=c(2,2))*

*arPosHi <- arima.sim(list(order=c(1,0,0), ar=0.9), n=100)*

*arPosLo <- arima.sim(list(order=c(1,0,0), ar=0.09), n=100)*

*arNegHi <- arima.sim(list(order=c(1,0,0), ar=-0.9), n=100)*

*arNegLo <- arima.sim(list(order=c(1,0,0), ar=-0.09), n=100)*

*periodogram(arPosHi); periodogram(arPosLo); periodogram(arNegHi);periodogram(arNegLo)*

*maPosHi <- arima.sim(list(order=c(0,0,1), ma=0.9), n=100)*

*maPosLo <- arima.sim(list(order=c(0,0,1), ma=0.09), n=100)*

*maNegHi <- arima.sim(list(order=c(0,0,1), ma=-0.9), n=100)*

*maNegLo <- arima.sim(list(order=c(0,0,1), ma=-0.09), n=100)*

*periodogram(maPosHi);periodogram(maPosLo);periodogram(maNegHi);periodogram(maNegLo)*

# Multiple Seasonality Modeling using TBATS

---

The TBATS model was introduced by De Livera, Hyndman & Snyder (2011, JASA). It is a generalization of the Holt-Winters model.

"TBATS" is an acronym denoting its salient features:

T for trigonometric regressors to model multiple-seasonalities

B for Box-Cox transformations

A for ARMA errors

T for trend

S for seasonality

The TBATS model is a generalization of the BATS model, which is similar except for lacking the trigonometric regressors.

The trigonometric output includes the periodicity and the number of harmonics/pairs for the time series.

The TBATS model can be fitted using the `tbats()` command in the forecast package for R.

# R code – Multiple Seasonality Modeling using TBATS

---

```
data(taylor)
```

```
# Taylor was defined taking the half hour electricity demand TS – msts is part of forecast pkg
```

```
# taylor <- msts(x, seasonal.periods=c(24 * 2, 24 * 2 * 7))
```

```
plot(taylor)
```

```
model <- tbats(taylor)
```

```
comp <- tbats.components(model)
```

```
plot(comp)
```

```
plot(forecast(model, h=100))
```

# Interpret TBATS Output in R

---

TBATS(0.999, {2,2}, 1, {<52.18,8>})\*

Box-Cox transformation of 0.999 (essentially doing nothing)

ARMA(2,2) errors,

Box-Cox damping parameter of 1 (doing nothing)

Seasonality modeled using 8 Fourier harmonics/pairs with period m=52.18

$$\begin{aligned}y_t &= \ell_{t-1} + b_{t-1} + s_{t-1} + \alpha d_t \\b_t &= b_{t-1} + \beta d_t \\s_t &= \sum_{j=1}^8 s_{j,t} \\s_{j,t} &= s_{j,t-1} \cos\left(\frac{2\pi jt}{52.18}\right) + s_{j,t-1}^* \sin\left(\frac{2\pi jt}{52.18}\right) + \gamma_1 d_t \\s_{j,t}^* &= -s_{j,t-1} \sin\left(\frac{2\pi jt}{52.18}\right) + s_{j,t-1}^* \cos\left(\frac{2\pi jt}{52.18}\right) + \gamma_2 d_t,\end{aligned}$$

where  $d_t$  is an ARMA(2,2) process and  $\alpha$ ,  $\beta$ ,  $\gamma_1$  and  $\gamma_2$  are smoothing parameters. Here the seasonality has been handled with 18 parameters (the sixteen initial values for  $s_{j,0}$  and  $s_{j,0}^*$  and the two smoothing parameters  $\gamma_1$  and  $\gamma_2$ ). The total number of degrees of freedom is 26 (the other 8 coming from the two smoothing parameters  $\alpha$  and  $\beta$ , the four ARMA parameters, and the initial level and slope values  $\ell_0$  and  $b_0$ ).

\* <https://robjhyndman.com/hyndsight/forecasting-weekly-data/>, Accessed May 2018

# Polynomial and Non-linear Regression

---

Polynomial regression


$$Y_t = \beta_0 + \beta_1 X_t + \beta_2 X_t^2 + u_t$$

The squared is represented directly in R using  $I(.)$

Non-linear regression

$$Y_t = a - be^{-cX_t}$$

# Survival Regression – Data censoring for time-to-event



# Survival Regression – Cox Regression

---

The Cox model is expressed by the hazard function denoted by  $h(t)$  that can be interpreted as the risk of dying at time  $t$  and estimated as follows:

$$h(t) = h_0(t) * \exp(b_1x_1 + b_2x_2 + \dots + b_px_p)$$

where,

$t$  represents the survival time

$h(t)$  is the hazard function determined by a set of  $p$  covariates the coefficients  $(x_1, x_2, \dots, x_p)$

measure the impact (i.e., the effect size) of covariates

$h_0(t)$  is the baseline hazard. It corresponds to the value of the hazard if all the covariates are equal to zero (the quantity  $\exp(0)$  equals 1).

The 't' in  $h(t)$  reminds us that the hazard may vary over time.

<http://www.sthda.com/english/wiki/cox-proportional-hazards-model>



# R code – Survival Regression


---

```
library("survival", lib.loc=~R/win-library/3.6")  
data("lung")  
uni.cox <- coxph(Surv(time, status) ~ sex, data = lung)  
mv.cox <- coxph(Surv(time, status) ~ age + sex + ph.ecog, data = lung)  
  
summary(mv.cox)
```

# Neural Networks for Time Series

---

- Fully connected Neural Network with 1 hidden layer
- NNAR  $(3,1,2)_{12}$  packages: `forecast (nnetar)` and `nnfor (elm and mlp)`



<https://otexts.org/fpp2/nnetar.html#fig:nnet2>

# R code – Neural Network Forecasting (nnfor)

---

```
library("nnfor", lib.loc=~R/win-library/3.3")
```

```
?elm
```

```
fit <- elm(AirPassengers, hd=10)
```

```
print(fit)
```

```
plot(fit)
```

```
frc <- forecast(fit,h=36)
```

```
plot(frc)
```

```
?mlp
```

```
fit2 <- mlp(AirPassengers, hd = c(10,5))
```

two hidden layers, one with 10 and one with 5

<https://kourentzes.com/forecasting/2019/01/16/tutorial-for-the-nnfor-r-package/>

# Recurrent Neural Network (RNN)

---

- RNNs address the temporal relationship of their inputs by maintaining an internal state.
- RNNs are biased towards learning patterns which occur in temporal order – i.e. they are less prone to learning random correlations which do not occur in temporal order.
- In theory, RNNs are absolutely capable of handling “long-term dependencies.” But in practice, RNNs suffer from vanishing gradient problem.
- Long Short Term Memory networks – usually just called “LSTMs” – are a special kind of RNN, capable of learning long-term dependencies\*.

\* <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>. Accessed Nov, 2016.

# Textbook Chapters

---

- Materials covered available in book:
  - MJK: Chapter 7, TSA: Chapter 13
- R package nnfor: <https://github.com/trnnick/nnfor/blob/master/R/>